

Hidden Markov Models

Peng Chen

School of Management
University of Science and Technology of China

December 10, 2020

Table of Contents

1 Hidden Markov Models

2 Inference in HMMs

- The forwards algorithm
- The forwards-backwards algorithm
- The Viterbi algorithm
- Forwards filtering, backwards sampling

3 Learning for HMMs

Table of Contents

1 Hidden Markov Models


2 Inference in HMMs

- The forwards algorithm
- The forwards-backwards algorithm
- The Viterbi algorithm
- Forwards filtering, backwards sampling

3 Learning for HMMs

Introduction

- ▶ A hidden Markov model or HMM consists of a discrete-time, discrete-state Markov chain, with hidden states $z_t \in \{1, \dots, K\}$, plus an observation model $p(x_t|z_t)$.



Notation

- ▶ State Space $\{1, \dots, K\}$, Observation Space.
- ▶ State Sequence: z_t , Observation Sequence: $x_t, t = 1, \dots, T$.
- ▶ Parameter: $\theta = (\pi, \mathbf{A}, \mathbf{B})$, where $\pi(i) = p(z_1 = i)$ is the initial state distribution, $A(i, j) = p(z_t = j | z_{t-1} = i)$ is the transition matrix, and \mathbf{B} are the parameters of the class-conditional densities $p(x_t | z_t = j)$.

Example (Boxes and Balls Model)

- ▶ There are 4 boxes with each containing 10 balls (red and white).
- ▶ We pick a box randomly with equal probability, and then pick a ball from this box randomly.
- ▶ State Space: $\{1, 2, 3, 4\}$, Observation Space: $\{R, W\}$.
- ▶ Suppose we have the observation sequence: $x_{1:5} = \{R, R, W, W, R\}$.

Box	Red	White
1	5	5
2	3	7
3	6	4
4	8	2

Example (Boxes and Balls Model)

- ▶ Initial state distribution $\boldsymbol{\pi} = (0.25, 0.25, 0.25, 0.25)^T$.
- ▶ The transition rule can be denoted by the transition matrix \mathbf{A} :

$$\begin{bmatrix} 0 & 1 & 0 & 0 \\ 0.4 & 0 & 0.6 & 0 \\ 0 & 0.4 & 0 & 0.6 \\ 0 & 0 & 0.5 & 0.5 \end{bmatrix}$$

- ▶ The class-conditional distribution matrix \mathbf{B} is

$$\begin{bmatrix} 0.5 & 0.5 \\ 0.3 & 0.7 \\ 0.6 & 0.4 \\ 0.8 & 0.2 \end{bmatrix}$$

Model Assumptions

- ▶ The observations are conditionally independent given the states:

$$P_{z_{t+1}|z_{1:t},x_{1:t}} = P_{z_{t+1}|z_t}.$$

- ▶ The state process is Markov:

$$P_{z_{t+1}|z_{1:t},x_{1:t}} = P_{z_{t+1}|z_t}.$$

- ▶ (Properties)

The joint process is Markov:

$$P_{x_{t+1:T},z_{t+1:T}|x_{1:t},z_{1:t}} = P_{x_{t+1:T},z_{t+1:T}|x_t,z_t}.$$

Given z_t , x_t is conditionally independent of everything else:

$$P_{x_t|x_{1:t-1},x_{t+1:T},z_{1:T}} = P_{x_t|z_t}.$$

HMM

- ▶ The corresponding joint distribution has the form

$$\begin{aligned} p(z_{1:T}, x_{1:T}) &= p(z_{1:T})p(x_{1:T}|z_{1:T}) \\ &= \left[p(z_1) \prod_{t=2}^T p(z_t|z_{t-1}) \right] \left[\prod_{t=1}^T p(x_t|z_t) \right]. \end{aligned}$$

- ▶ $p(z_{1:T}) = p(z_1) \prod_{t=2}^T p(z_t|z_{t-1})$ (Bayes rule & The state is Markov)

- ▶ $p(x_{1:T}|z_{1:T}) = \prod_{t=1}^T p(x_t|z_t)$
(The observations are conditionally independent given the states)

Applications

- ▶ HMMs have the advantage over Markov models in that they can represent long-range dependencies between observations, mediated via the latent variables.
- ▶ Two common scenarios:
 - Online scenario: compute $p(z_t|x_{1:t})$.
 - Offline scenario: compute $p(z_t|x_{1:T})$.

Examples of Applications

► Automatic speech recognition:

Here x_t represents features extracted from the speech signal, and z_t represents the word that is being spoken. The transition model $p(z_t|z_{t-1})$ represents the language model, and the observation model $p(x_t|z_t)$ represents the acoustic model.

► Activity recognition

Here x_t represents features extracted from a video frame, and z_t is the class of activity.

Examples of Applications

- ▶ Part of speech tagging

Here x_t represents a word, and z_t represents its part of speech (noun, verb, adjective, etc.)

- ▶ Gene finding

Here x_t represents the DNA nucleotides (A,C,G,T), and z_t represents whether we are inside a gene-coding region or not.

- ▶ Protein sequence alignment

Here x_t represents an amino acid, and z_t represents whether this matches the latent consensus sequence at this location. This model is called a profile HMM.

Table of Contents

1 Hidden Markov Models

2 Inference in HMMs

- The forwards algorithm
- The forwards-backwards algorithm
- The Viterbi algorithm
- Forwards filtering, backwards sampling

3 Learning for HMMs

Types of inference problems

- ▶ We discuss how to infer the hidden state sequence of an HMM, assuming the parameters are known.
- ▶ Consider an example called the occasionally dishonest casino.
- ▶ In this model, $x_t \in \{1, 2, \dots, 6\}$ represents which dice face shows up, and $z_t \in \{1(\text{fair}), 2(\text{loaded})\}$ represents the identity of the dice that is being used.




Figure 17.9 An HMM for the occasionally dishonest casino. The blue arrows visualize the state transition diagram **A**. Based on (Durbin et al. 1998, p54).

Types of inference problems

- ▶ We can just see the rolls and want to infer which dice is being used. Following are different kinds of inference.
- ▶ **Filtering** means to compute the belief state $p(z_t|x_{1:t})$ online, or recursively, as the data streams in.
- ▶ **Smoothing** means to compute $p(z_t|x_{1:T})$ offline, given all evidence.

Listing 17.1 Example output of casinoDemo

```
Rolls: 664153216162115234653214356634261655234232315142464156663246
Die: LLLLLLLLLLLLLLFFFFFFFFLLLLLLLLLLLLLLLLFFFFFFFFFFFFFFFFFFFFFFFFLLLLLLLL
```

Types of inference problems

- ▶ **Fixed lag smoothing** is an compromise between online and offline estimation for its computation of $p(z_{t-l}|x_{1:t})$, where $l > 0$ is called the lag.
- ▶ **Prediction** We might want to predict the future given the past, i.e., to compute $p(z_{t+h}|x_{1:t})$, where $h > 0$ is called the prediction horizon.
- ▶ **MAP estimation** This means computing $\operatorname{argmax}_{z_{1:T}} p(z_{1:T}|x_{1:T})$, which is a most probable state sequence (Viterbi decoding).

Types of inference problems

- ▶ **Posterior samples** We can obtain more information from the sample paths sampled from the posterior, $z_{1:T} \sim p(z_{1:T}|x_{1:T})$, than the sequence of marginals computed by smoothing.
- ▶ **Probability of the evidence** We can compute the probability of the evidence, $p(x_{1:T})$, by summing up over all hidden paths, $p(x_{1:T}) = \sum_{z_{1:T}} p(z_{1:T}, x_{1:T})$.

Types of inference problems




Figure 17.11 The main kinds of inference for state-space models. The shaded region is the interval for which we have data. The arrow represents the time step at which we want to perform inference. t is the current time, T is the sequence length, ℓ is the lag and h is the prediction horizon. See text for details.

The forwards algorithm

- ▶ To compute the filtered marginals, $\alpha_t = p(z_t|x_{1:t})$ in an HMM.
- ▶ The algorithm has two steps: the prediction step and the update step.
- ▶ **Prediction step** We compute the one-step-ahead predictive density,

$$p(z_t = j|x_{1:t-1}) = \sum_i p(z_t = j|z_{t-1} = i)p(z_{t-1} = i|x_{1:t-1}),$$

which acts as the new prior for time t .

The forwards algorithm

- **Update step** We absorb the observed data from time t using Bayes rule,

$$\begin{aligned}
 \alpha_t(j) &\triangleq p(z_t = j | x_{1:t}) = p(z_t = j | x_t, x_{1:t-1}) \\
 &= \frac{p(x_t | z_t = j, x_{1:t-1}) p(z_t = j | x_{1:t-1})}{p(x_t | x_{1:t-1})} \\
 &= \frac{p(x_t | z_t = j) p(z_t = j | x_{1:t-1})}{Z_t} \\
 &\propto p(x_t | z_t = j) p(z_t = j | x_{1:t-1}),
 \end{aligned}$$

where

$$\begin{aligned}
 Z_t &= \sum_j p(x_t | z_t = j) p(z_t = j | x_{1:t-1}) \\
 &= \sum_j p(x_t, z_t = j | x_{1:t-1}) = p(x_t | x_{1:t-1}).
 \end{aligned}$$

The forwards algorithm

- ▶ The distribution $p(z_t|x_{1:t})$ is called the (filtered) belief state at time t , and is a vector of K numbers, denoted by α_t .
- ▶ We can rewrite the update in the matrix-vector form:

$$\alpha_t \propto \phi_t \odot (\Psi^T \alpha_{t-1}),$$

where $\phi_t(j) = p(x_t|z_t = j)$ is the local evidence at time t , $\Psi(i, j) = p(z_t = j|z_{t-1} = i)$ is the transition matrix, and \odot is the Hadamard product.

The forwards algorithm

- ▶ In addition to computing the hidden states, we can use this algorithm to compute the log probability of the evidence:

$$\log p(x_{1:T}|\boldsymbol{\theta}) = \sum_{t=1}^T \log p(x_t|x_{1:t-1}) = \sum_{t=1}^T \log Z_t.$$

Algorithm 17.1: Forwards algorithm

- 1 Input: Transition matrices $\psi(i, j) = p(z_t = j | z_{t-1} = i)$, local evidence vectors $\psi_t(j) = p(\mathbf{x}_t | z_t = j)$, initial state distribution $\pi(j) = p(z_1 = j)$;
 - 2 $[\alpha_1, Z_1] = \text{normalize}(\psi_1 \odot \boldsymbol{\pi})$;
 - 3 **for** $t = 2 : T$ **do**
 - 4 $[\alpha_t, Z_t] = \text{normalize}(\psi_t \odot (\boldsymbol{\Psi}^T \alpha_{t-1}))$;
 - 5 Return $\alpha_{1:T}$ and $\log p(\mathbf{y}_{1:T}) = \sum_t \log Z_t$;
 - 6 Subroutine: $[\mathbf{v}, Z] = \text{normalize}(\mathbf{u}) : Z = \sum_j u_j; \quad v_j = u_j/Z$;
-

Examples

- ▶ Consider the Boxes and Balls Model with observation sequence $\{R, R, W, W, R\}$ and parameters in page 7, and then we have

$$\alpha_1 = (0.227, 0.136, 0.273, 0.364)$$

$$\alpha_2 = (0.121, 0.136, 0.291, 0.452)$$

$$\alpha_3 = (0.158, 0.364, 0.304, 0.173)$$

$$\alpha_4 = (0.363, 0.343, 0.199, 0.095)$$

$$\alpha_5 = (0.353, 0.164, 0.240, 0.243)$$

$$p(x_{1:5}|\theta) = 0.032$$

The forwards-backwards algorithm

- ▶ To compute the smoothed marginals $p(z_t = j | x_{1:T})$.
- ▶ We break the chain into two parts, the past and the future, by conditioning on z_t :

$$p(z_t = j | x_{1:T}) \propto p(z_t = j, x_{t+1:T} | x_{1:t}) \propto p(z_t = j | x_{1:t}) p(x_{t+1:T} | z_t = j).$$

- ▶ Define $\beta_t(j) \triangleq p(x_{t+1:T} | z_t = j)$ as the conditional likelihood of future evidence given that the hidden state at time t is j .
- ▶ Define $\gamma_t(j) \triangleq p(z_t = j | x_{1:T})$ as the desired smoothed posterior marginal. Thus, $\gamma_t(j) \propto \alpha_t(j) \beta_t(j)$.

The forwards-backwards algorithm

- ▶ We now describe how to recursively compute the β 's in a right-to-left fashion.
- ▶ If we have already computed β_t , we can compute β_{t-1} as follows:

$$\begin{aligned}\beta_{t-1}(i) &= p(x_{t:T} | z_{t-1} = i) \\ &= \sum_j p(z_t = j, x_t, x_{t+1:T} | z_{t-1} = i) \\ &= \sum_j p(x_{t+1:T} | z_t = j, z_{t-1} = i, x_t) p(z_t = j, x_t | z_{t-1} = i) \\ &= \sum_j p(x_{t+1:T} | z_t = j) p(x_t | z_t = j, z_{t-1} = i) p(z_t = j | z_{t-1} = i) \\ &= \sum_j \beta_t(j) \phi_t(j) \Psi(i, j).\end{aligned}$$

The forwards-backwards algorithm

- ▶ We rewrite the equation in matrix-vector form as

$$\beta_{t-1} = \Psi(\phi_t \odot \beta_t).$$

- ▶ The base case is

$$\beta_T(i) = p(x_{T+1:T} | z_T = i) = p(\emptyset | z_T = i) = 1,$$

which is the probability of a non-event.

Two-slice smoothed marginals

- ▶ In order to estimate the parameters of the transition matrix using EM, we need to compute:

$$N_{ij} = \sum_{t=1}^{T-1} E[I(z_t = i, z_{t+1} = j) | x_{1:T}] = \sum_{t=1}^{T-1} p(z_t = i, z_{t+1} = j | x_{1:T}).$$

Two-slice smoothed marginals

- ▶ The term $p(z_t = i, z_{t+1} = j | x_{1:T})$ is called a (smoothed) two-slice marginal, and can be computed as follows:

$$\begin{aligned}
 \xi_{t,t+1}(i, j) &\triangleq p(z_t = i, z_{t+1} = j | x_{1:T}) \\
 &= p(z_t | x_{1:T}) p(z_{t+1} | z_t, x_{1:T}) \\
 &\propto p(z_t | x_{1:t}) p(x_{t+1:T} | z_t) p(z_{t+1} | z_t, x_{1:T}) \\
 &= p(z_t | x_{1:t}) p(x_{t+1:T} | z_t) p(z_{t+1} | z_t, x_{t+1:T}) \\
 &= p(z_t | x_{1:t}) p(x_{t+1:T} | z_t, z_{t+1}) p(z_{t+1} | z_t) \\
 &\propto p(z_t | x_{1:t}) p(x_{t+1} | z_{t+1}) p(x_{t+2:T} | z_{t+1}) p(z_{t+1} | z_t) \\
 &= \alpha_t(i) \phi_{t+1}(j) \beta_{t+1}(j) \psi(i, j)
 \end{aligned}$$

- ▶ In matrix-vector form, we have

$$\xi_{t,t+1} \propto \Psi \odot (\alpha_t(\phi_{t+1} \odot \beta_{t+1})^T).$$

Time and space complexity

- ▶ A straightforward implementation of FB takes $O(K^2T)$ time, since we must perform a $K \times K$ matrix multiplication at each step.
- ▶ In some cases, the bottleneck is memory, not time. It is possible to devise a simple divide-and-conquer algorithm that reduces the space complexity from $O(KT)$ to $O(K\log T)$ at the cost of increasing the running time from $O(K^2T)$ to $O(K^2T\log T)$.

The Viterbi algorithm

- ▶ To compute $z^* = \operatorname{argmax}_{z_{1:T}} p(z_{1:T} | x_{1:T})$.
- ▶ It is equivalent to computing a shortest path through the trellis diagram in Figure 17.12 with the weight of a path z_1, z_2, \dots, z_T given by the log probability $\log p(z_{1:T}, x_{1:T})$




Figure 17.12 The trellis of states vs time for a Markov chain. Based on (Rabiner 1989).

MAP vs MPE

- ▶ The **jointly** most probable sequence of states is not necessarily the same as the sequence of **marginally** most probable states.
- ▶ The former is what Viterbi computes.
- ▶ The latter is given by the maximizer of the posterior marginals or MPM:

$$\hat{z} = (\operatorname{argmax}_{z_1} p(z_1 | x_{1:T}), \dots, \operatorname{argmax}_{z_T} p(z_T | x_{1:T})).$$

	$X_1 = 0$	$X_1 = 1$	
$X_2 = 0$	0.04	0.3	0.34
$X_2 = 1$	0.36	0.3	0.66
	0.4	0.6	

The Viterbi algorithm

- ▶ Define the probability of ending up in state j at time t , given that we take the most probable path, as

$$\delta_t(j) \triangleq \max_{z_1, \dots, z_{t-1}} p(z_{1:t-1}, z_t = j | x_{1:t}).$$

- ▶ The most probable path to state j at time t must consist of the most probable path to some other state i at time $t-1$, followed by a transition from i to j . Hence,

$$\delta_t(j) = \max_i \delta_{t-1}(i) \psi(i, j) \phi_t(j).$$

The Viterbi algorithm

- ▶ Let

$$a_t(j) = \underset{i}{\operatorname{argmax}} \delta_{t-1}(i) \psi(i, j) \phi_t(j),$$

and it is the most likely previous state on the most probable path to $z_t = j$.

- ▶ We initialize by setting $\delta_1(j) = \pi_j \phi_1(j)$ and terminate by computing the most probable final state

$$z_T^* = \underset{i}{\operatorname{argmax}} \delta_T(i).$$

- ▶ We can then compute the most probable sequence of states using traceback:


$$z_t^* = a_{t+1}(z_{t+1}^*).$$

The Viterbi algorithm


- ▶ In order to avoid numerical underflow, we can normalize the δ_t terms at each step.
- ▶ Unlike the forwards-backwards case, we can easily work in the log domain ($\log \max = \max \log, \log \sum \neq \sum \log$), which can result in a significant speedup in the case of Gaussian observation models.

Example

- Consider a simple HMM with observation space $\{C_1, C_2, \dots, C_7\}$, and its transition matrix and class-conditional probability is expressed in the following figure.
- Suppose we observe the sequence of observations $\{C_1, C_3, C_4, C_6\}$.
- The model starts in state S1.



(a)



(b)

Example

$$t = 1, \quad \delta_1(1) = 0.5,$$

$$\delta_1(2) = 0,$$

$$\delta_1(3) = 0,$$

$$t = 2, \quad \delta_2(1) = \delta_1(1)\psi(1, 1)\phi_1(C_3) = 0.5 \times 0.3 \times 0.3 = 0.045,$$

$$\delta_2(2) = \delta_1(1)\psi(1, 2)\phi_2(C_3) = 0.5 \times 0.7 \times 0.2 = 0.07,$$

$$\delta_2(3) = 0;$$

$$t = 3, \quad \delta_3(1) = 0,$$

$$\begin{aligned} \delta_3(2) &= \max\{\delta_2(1)\psi(1, 2)\phi_2(C_4), \delta_2(2)\psi(2, 2)\phi_2(C_4)\} \\ &= \max\{0.02205, 0.0441\} = 0.0441, \end{aligned}$$

$$\delta_3(3) = \delta_2(1)\psi(2, 3)\phi_3(C_4) = 0.07 \times 0.1 \times 0.1 = 0.0007;$$

Example


► (continue)

$$t = 4, \quad \delta_4(1) = 0,$$

$$\delta_4(2) = 0,$$

$$\begin{aligned} \delta_4(3) &= \max\{\delta_3(2)\psi(2, 3)\phi_3(C_6), \delta_3(3)\psi(3, 3)\phi_3(C_6)\} \\ &= \max\{0.0022, 0.0001\} = 0.0022; \end{aligned}$$

- So, we have $z_4^* = S_3$ and then we can calculate that $z_3^* = S_2, z_2^* = S_2, z_1^* = S_1$ using traceback.



The Viterbi algorithm

- ▶ The time complexity of Viterbi is clearly $O(K^2T)$ in general, and the space complexity is $O(KT)$, both the same as forwards-backwards.
- ▶ The Viterbi algorithm can be extended to return the top N paths, which is called the N-best list.

Forwards filtering, backwards sampling

- ▶ To sample paths from the posterior: $z_{1:T}^s \sim p(z_{1:T}|x_{1:T})$.
- ▶ One way is to do as follow:
run forwards backwards, to compute the two-slice smoothed posteriors, $p(z_{t-1,t}|x_{1:T})$; next compute the conditionals $p(z_t|z_{t-1}, x_{1:T})$ by normalizing; sample from the initial pair of states, $z_{1,2}^* \sim p(z_{1,2}|x_{1:T})$; finally, recursively sample $z_t^* \sim p(z_t|z_{t-1}^*, x_{1:T})$.
(a forwards-backwards pass and an additional forwards sampling pass.)

Forwards filtering, backwards sampling

- ▶ An alternative is to do the forwards pass, and then perform sampling in the backwards pass.
- ▶ We write the joint from right to left :

$$p(z_{1:T}|x_{1:T}) = p(z_T|x_{1:T}) \prod_{t=T-1}^1 p(z_t|z_{t+1}, x_{1:T}).$$

- ▶ We can then sample z_t given future sampled states using

$$z_t^s \sim p(z_t|z_{t+1:T}, x_{1:T}) = p(z_t|z_{t+1}, z_{t+2:T}, x_{1:t}, x_{t+1:T}) = p(z_t|z_{t+1}^s, x_{1:t})$$

Forwards filtering, backwards sampling

- ▶ The sampling distribution is given by

$$\begin{aligned}
 p(z_t = i | z_{t+1} = j, x_{1:t}) &= p(z_t | z_{t+1}, x_{1:t}, x_{t+1}) \\
 &= \frac{p(z_{t+1}, z_t | x_{1:t+1})}{p(z_{t+1} | x_{1:t+1})} \\
 &\propto \frac{p(x_{t+1} | z_{t+1}, z_t, x_{1:t}) p(z_{t+1}, z_t | x_{1:t})}{p(z_{t+1} | x_{1:t+1})} \\
 &= \frac{p(x_{t+1} | z_{t+1}) p(z_{t+1} | z_t, x_{1:t}) p(z_t | x_{1:t})}{p(z_{t+1} | x_{1:t+1})} \\
 &= \frac{\phi_{t+1}(j) \psi(i, j) \alpha_t(i)}{\alpha_{t+1}(j)}
 \end{aligned}$$

- ▶ The base case is $z_T^s \sim p(z_T = i | x_{1:T}) = \alpha_T(i)$.

Table of Contents

1 Hidden Markov Models

2 Inference in HMMs

- The forwards algorithm
- The forwards-backwards algorithm
- The Viterbi algorithm
- Forwards filtering, backwards sampling

3 Learning for HMMs

Learning for HMMs

- ▶ To estimate the parameters $\theta = (\pi, A, B)$, where $\pi(i) = p(z_1 = i)$ is the initial state distribution, $A(i, j) = p(z_t = j | z_{t-1} = i)$ is the transition matrix, and B are the parameters of the class-conditional densities $p(x_t | z_t = j)$.
- ▶ Case1: $z_{1:T}$ is **observed** We can easily compute the MLEs for θ .
- ▶ Case2: $z_{1:T}$ is **hidden** We can estimate θ with the EM (Baum-Welch) algorithm.

E step

- Given the sample $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, the expected complete data log likelihood is

$$Q(\boldsymbol{\theta}, \boldsymbol{\theta}^{old}) = \sum_{k=1}^K E[N_k^1] \log \pi_k + \sum_{j=1}^K \sum_{k=1}^K E[N_{jk}] \log A_{jk} \\ + \sum_{i=1}^N \sum_{t=1}^{T_i} \sum_{k=1}^K p(z_t = k | \mathbf{x}_i, \boldsymbol{\theta}^{old}) \log p(x_{i,t} | \phi_k),$$

E step

$$\blacktriangleright E[N_k^1] = \sum_{i=1}^N p(z_{i1} = k | \mathbf{x}_i, \boldsymbol{\theta}^{old}) = \sum_{i=1}^N \gamma_{i,k}(1)$$

(the number in state k at time 1)

$$\blacktriangleright E[N_{jk}] = \sum_{i=1}^N \sum_{t=2}^{T_i} p(z_{i,t-1} = j, z_{i,t} = k | \mathbf{x}_i, \boldsymbol{\theta}^{old}) = \sum_{i=1}^N \sum_{t=2}^{T_i} \xi_{i,j,k}(t)$$

(the number of transitions from state j to state k)

$$\blacktriangleright E[N_j] = \sum_{i=1}^N \sum_{t=1}^{T_i} p(z_{i,t} = j | \mathbf{x}_i, \boldsymbol{\theta}^{old}) = \sum_{i=1}^N \sum_{t=1}^{T_i} \gamma_{i,t}(j)$$

(the total number of times in state j)

M step

- ▶ We have that the M step for \mathbf{A} and $\boldsymbol{\pi}$ is to just normalize the expected number:

$$\hat{A}_{jk} = \frac{E[N_{jk}]}{\sum_{k'} E[N_{jk'}]}, \hat{\pi}_k = \frac{E[N_k^1]}{N}.$$

- ▶ $\hat{B}_{jl} = \frac{E[M_{jl}]}{E[N_j]}$, where $E[M_{jl}] = \sum_{i=1}^N \sum_{t=1}^{T_i} \gamma_{i,t}(j) I(x_{i,t} = l)$.

Initialization

- ▶ Use some fully labeled data to initialize the parameters.
- ▶ Initially ignore the Markov dependencies, and estimate the observation parameters using the standard mixture model estimation methods, such as K-means or EM.
- ▶ Randomly initialize the parameters, use multiple restarts, and pick the best solution.

Model selection

- ▶ Two main issues: how many states, and what topology to use for the state transition diagram.
- ▶ Choosing the number of hidden states
 - (i) Use grid-search over a range of K 's, using as an objective function cross-validated likelihood, the BIC score, or a variational lower bound to the log-marginal likelihood.
 - (ii) Use reversible jump MCMC.
 - (iii) Use variational Bayes to “extinguish” unwanted components.
 - (iv) Use an “infinite HMM”, which is based on the hierarchical Dirichlet process.

Model selection

- ▶ Structure learning
- ▶ To learn the structure of the state transition diagram, not the structure of the graphical model (which is fixed).
- ▶ Alternatively, one can pose the problem as MAP estimation using a minimum entropy prior.